Sri Omkar D

sri.omkar.d@gmail.com | **+1 951-545-2146** | **sriomkar.com** | **linkedin.com/in/sri-omkar-58r4r4r8**

Professional Summary

Senior Data Analyst & Analytics Engineer with 7+ years using Python, SQL, PySpark, Databricks, and Tableau to build reliable data pipelines and executive reporting. Proven track record of automating manual business processes, diagnosing root causes behind metric anomalies, and cutting batch runtimes in half to protect leadership reporting SLAs. Strong background in partnering with business stakeholders to turn complex data into clear dashboards and revenue-impacting decisions.

Technical Skills

- **Languages & Scripting**: SQL (T-SQL, PL/SQL, Spark SQL), Python, PySpark
- **Analytics & BI**: Tableau, Power BI, KPI reporting, root-cause analysis, variance analysis, forecasting datasets
- **Data Platforms & Storage**: Databricks, Snowflake, Azure (ADLS Gen2), AWS (S3, Redshift, Glue, Athena), SQL Server, Oracle
- **Orchestration & DevOps**: Airflow, Git, GitHub Actions, Docker, CI/CD, scheduled alerting
- **Data Quality & Modeling**: Source-to-target reconciliation, dimensional modeling (star/snowflake, SCD Type 2), automated data archiving, schema enforcement
- **ETL Tools**: IBM DataStage, Microsoft SSIS, Ab Initio, AWS Glue

Nortek Consulting INC (Optum) **Florida**
Data Engineer **Oct 2025 - Present**

- Own, monitor, and optimize **Databricks** and **PySpark** data pipelines feeding daily business and KPI reporting layers.
- Cut nightly load times by 50% (from 7:30 AM down to 3:00 AM) by re-partitioning tables and tuning query logic, ensuring leadership dashboards are refreshed hours before the business day starts.
- Design and maintain Tableau dashboards and **SQL** reporting views, giving stakeholders direct visibility into operational metrics, performance trends, and volume variances.
- Automate manual data processes using **Python** and PySpark, including automated data archiving, dynamic schema validation for nested inputs, and automated failure alerts that give business teams hours back each week.
- Tuned Delta Lake tables using OPTIMIZE and Z-Ordering on high-cardinality join keys (date, provider ID, transaction ID), cutting Tableau workbook query response times from over 40 seconds to under 4 seconds.
- Lead root-cause analysis (**RCA**) when metric trends or source feeds show unexpected anomalies; built automated source-to-target reconciliation checks to catch issues before numbers hit downstream reports.
- Partner directly with business teams to clarify reporting requirements, explain the data behind key outcomes, and mentor junior analysts on SQL tuning, PySpark best practices, and runbooks.
- Partnered directly with leadership on high-priority, ad-hoc deep dives using complex SQL (CTEs, window functions) to explain unexpected margin dips and regional performance drops.

Akkodis (Herc Rentals Inc.) **Florida**
Data Engineer **Nov 2024 – Sep 2025**

- Developed **Python**-based ingestion workflows to route high-volume operational data from **RentalMan**, **AS400**, **Oracle OLTP**, and **Teradata-based systems** into an **AWS S3** data lake for rental, inventory, billing, and branch operations reporting.
- Built near-real-time equipment monitoring pipelines using **AWS Kinesis** and **Databricks Streaming** to process semi-structured **JSON/XML** telemetry payloads, including **GPS location**, **engine metrics**, **fuel usage**, and fleet activity for maintenance and utilization reporting.
- Created and maintained **PySpark** and **Spark SQL** pipelines in **Databricks** to clean, normalize, and join datasets across rental contracts, fleet inventory, billing, customer operations, and equipment usage.
- Reduced daily batch runtimes by approximately **25–30%** by tuning **PySpark** partition strategies, optimizing **Spark SQL** joins, and replacing unnecessary full-refresh cycles with scheduled incremental load patterns.
- Decreased pipeline rerun effort by approximately **30–40%** by breaking long **Databricks** transformation chains into modular, restartable stages with intermediate curated outputs for rental, billing, and fleet reporting loads.
- Added **data validation** and **source-to-target reconciliation** checks across pipeline stages, including record counts, duplicate checks, null handling, and source-to-target comparisons to improve data reliability before business reporting.
- Built SQL-based curated views and reporting layers using **Teradata-based datasets**, **fact/dimension modeling**, **materialized views**, and partition-aware **SQL** design to support **Qlik Sense** dashboards for fleet utilization, rental performance, and financial reporting.
- Orchestrated scheduled workflows using **Apache Airflow** and **SQL Server Agent**, including dependency handling, retry logic, monitoring, and troubleshooting for failed loads.

Quess Corp (Blue Yonder India Pvt Ltd.) **India**
Data Engineer **June 2022 - July 2023**

- Owned **ETL** for retail supply chain data. Inventory, POS, orders, product movement. Sources were **ERP** systems, flat files, APIs. Target was the **Azure** data lake and warehouse.
- Every transformation started with a source-to-target mapping. What comes in, what goes out, the business rules in between. Nothing got built until the mapping was signed off.
- Wrote the transformation logic in **PySpark** and **SQL**. Merged scattered enrichment steps into single unified runs. Cut batch runtimes by **25–30%.**
- Tested transformations before they touched production. Ran them on sample data, checked outputs row by row, then pointed them at the real tables.
- Built validation into every stage. Duplicate checks, null handling, source-to-target row comparisons. A load that failed a check stopped there instead of corrupting downstream.

- Tuned the slow parts. Profiled queries on **Snowflake** and Exasol, fixed joins, redrew partitions on the big tables.
- Ran the batches in **Airflow** and owned what happened after. DAGs, retries, SLA alerts. When something broke overnight, I found why and fixed it. Wrote the docs so the next person wouldn't have to guess.

Trigent Software Pvt Ltd. (Accenture Solutions Pvt Ltd.) **India**
Data Engineer **Jan 2020 - June 2022**

- Migrated legacy **Ab Initio** workflows to **AWS**. Mapped every old job first, then rebuilt from the map.
- Landed data in **S3** across raw, cleansed, and curated zones. No mixing.
- Rebuilt transforms in **PySpark** on **Glue**. Registered every table in the **Glue Data Catalog**.
- Modeled **Redshift** dims and facts. Set dist and sort keys for how reports actually query.
- Ran old and new side by side before cutover. Row counts, column parity. Switched over only when the numbers tied.
- Cut runtime **35–40%** on financial pipelines. Diagnosed skew in **Spark UI**, applied salting, broadcast joins, **AQE**.
- Baked source-to-target validation into every job. Regulatory reporting, zero room for wrong numbers.
- **Airflow** for scheduling, **CI/CD** with Git and Jenkins across DEV, QA, PROD. Handed over with runbooks.
- Automated the pipeline plumbing in **Python**. File checks, rerun scripts, alerts between jobs. The unglamorous code that keeps batches moving.
- Used **Athena** to query raw **S3** data directly during migration. Fastest way to compare legacy output against new tables without loading anything.

Experis IT (Thomson Reuters) **India**
ETL Developer **June 2017 - Sep 2019**

- Built **SSIS** packages to extract from legacy **SQL Server** databases into an **S3** data lake. Control flow for orchestration, data flow for transforms. Sales, subscription, revenue data.
- Where SSIS wasn't the right fit, built the transforms in Glue, **PySpark**, and HiveQL.
- Wrote heavy **SQL** daily. Window functions, subqueries, multi-table joins. For validation, ad-hoc analysis, tuning.
- Ran data checks every morning. Missing records, duplicates, source-to-target counts, revenue variance. Before the reports went out.
- Sped up slow batches **20–25%. Parquet** outputs, fixed partitioning.
- Scheduled and watched everything on **SQL Server Agent**. Diagnosed failed loads, protected **SLAs**.
- Wrote **T-SQL** procedures, functions, and views behind the ETL and reports. Automated routine ops with **UNIX** scripts.

Education:

Master's in Information Technology Management **2024**
California Baptist University**,** Riverside, CA

B. Tech in Computer Science and Engineering **2017**
JNTU-K University, India