

Sri Omkar D

sri.omkar.d@gmail.com | +1 951-545-2146 | sriomkar.com | linkedin.com/in/sri-omkar-58r4r4r8

PROFESSIONAL SUMMARY

- **ETL / Data Engineer with 7+ years of experience** designing, developing, and optimizing enterprise data integration pipelines, including 5+ years of hands-on ETL development across cloud and on-premise platforms.
- Strong working knowledge of **IBM DataStage** (parallel jobs, server jobs, sequence jobs) and **Microsoft SSIS** (control flow, data flow, package deployment) — experienced in analyzing, troubleshooting, and migrating legacy DataStage/SSIS workloads to modern cloud architectures.
- Proven experience **leading and supporting cloud migration and ETL modernization initiatives**: reverse-engineering legacy jobs, mapping dependencies and transformations, rebuilding logic in AWS-native services (S3, Glue, Redshift, Lambda), and validating parity through parallel runs and reconciliation.
- Deep expertise in **SQL (T-SQL, PL/SQL, Spark SQL)** — complex joins, window functions, stored procedures, data mapping, and performance tuning across relational databases (SQL Server, Oracle, Teradata, DB2) and enterprise data warehouses (Redshift, Snowflake).
- Rigorous **data quality** discipline: validation, source-to-target reconciliation, parity checks, regression testing, and UAT for migrated pipelines.
- **Production support** veteran: batch monitoring, job scheduling (SQL Server Agent, Airflow, Control-M concepts), root-cause analysis, and hypercare during cutovers.
- **Healthcare data governance** experience: HIPAA/HITRUST-aligned engineering — PHI data masking, RBAC, audit-ready lineage, and minimum-necessary handling inside ETL workflows.
- Collaborative partner to data architects, cloud teams, DBAs, and business stakeholders; Agile/Scrum, CI/CD, and thorough technical documentation.

TECHNICAL SKILLS

ETL Tools: IBM DataStage (Parallel / Server / Sequence Jobs), Microsoft SSIS (Control Flow, Data Flow), Ab Initio, AWS Glue

Cloud Platforms: AWS (S3, Glue, Glue Data Catalog, Redshift, Lambda, Athena, EMR, Kinesis, IAM, CloudWatch), Microsoft Azure (ADLS Gen2, Azure Databricks, Key Vault, Monitor)

Programming: SQL (T-SQL, PL/SQL, Spark SQL, HiveQL), Python, PySpark

Databases & Warehousing: SQL Server, Oracle, Teradata, DB2, Snowflake, Amazon Redshift, Exasol; dimensional modeling (star/snowflake schema, SCD Type 2), MPP architectures

Scheduling & Orchestration: SQL Server Agent, Apache Airflow, Control-M concepts, cron

Data Quality & Governance: Source-to-target reconciliation, parity checks, data profiling, cleansing, de-duplication, data lineage, HIPAA, HITRUST CSF, data masking, RBAC

DevOps: Git, GitHub Actions, Jenkins, Terraform, Docker, CI/CD

BI & Reporting: Tableau, Power BI, Qlik Sense

PROFESSIONAL EXPERIENCE

Optum — Florida

Data Engineer | Nov 2024 – Present

- Analyze and **reverse-engineer legacy IBM DataStage and SSIS jobs** — mapping dependencies, hardcoded business logic, and source-to-target transformations — to assess migration readiness for cloud modernization.
- **Rebuild legacy ETL logic as AWS-native pipelines** (S3, Glue, Lambda, Redshift) with cloud-native logging, scheduling, and error handling; decommission legacy jobs after parity sign-off.

- Engineer **HIPAA/HITRUST-aligned data pipelines**: dynamic data masking and RBAC for PHI/ePHI, minimum-necessary field handling, and end-to-end automated lineage from raw sources to reporting layers for audit readiness.
- Perform rigorous **data validation and reconciliation** — row counts, parity checks, schema validation, regression testing, and UAT — running modernized pipelines in parallel with legacy jobs for column-by-column verification.
- **Tune ETL performance**: refactored inefficient legacy SQL, optimized partitioning/clustering/indexing for the enterprise data warehouse, shrinking batch execution windows to meet SLAs.
- Provide **production support** for DataStage/SSIS batch jobs: overnight monitoring, **break-fix changes**, root-cause analysis on failures (e.g., source file format changes), and data-quality anomaly triage before downstream claims/reporting impact; implement **preventive fixes** to eliminate recurring failures.
- Partner with **data architects, DBA teams, and business stakeholders** on migration planning, cutover execution, and hypercare stabilization; author source-to-target mapping specifications and operational runbooks.
- Plan and execute **trial migrations and cutover**: dual-parallel validation (row-by-row, column-by-column parity) before go-live, with **hypercare support** through post-migration stabilization.
- Drive **data profiling, cleansing, and de-duplication** ahead of migration; define archival strategies for legacy data not carried into the cloud platform.
- Establish **operational standards for platform validation and data integrity** across the data lifecycle; identify and mitigate single points of failure (SPOF) to harden pipeline resiliency.

Quess Corp (Blue Yonder India Pvt Ltd.) — India

Data Engineer | June 2022 – July 2023

- Developed and maintained **ETL workflows** integrating retail supply chain data (inventory, POS, orders, product movement) from ERP systems, flat files, and APIs into Azure-based data lake and warehouse layers.
- Wrote **complex data mappings and transformation logic** (PySpark, SQL), consolidating iterative enrichment steps into unified operations that cut batch transformation runtimes by ~25–30%.
- Performed **SQL tuning and query profiling** across Snowflake and Exasol; optimized joins and partition boundaries for large-scale datasets.
- Implemented **data validation checks** — duplicate detection, null handling, source-to-target comparisons — improving reliability of downstream analytics.
- Administered **scheduled batch workflows (Apache Airflow)**: DAG dependencies, retries, SLA monitoring, and alerting; produced technical documentation for data flows and mappings.
- Modeled **dimensional warehouse structures** (star schema, SCD Type 2) to support supply chain analytics and reporting workloads.
- Partnered directly with **business users and data management teams** to translate business requirements into technical ETL solutions.

Trigent Software Pvt Ltd. (Accenture Solutions Pvt Ltd.) — India

Data Engineer | Jan 2020 – June 2022

- **Led migration of legacy Ab Initio ETL workflows to AWS cloud-native architecture** (S3 data lake across raw/cleansed/curated layers, Glue Data Catalog, Redshift) — the same analyze → translate → re-platform motion as DataStage/SSIS modernization.
- Reduced runtime by ~35–40% on high-volume financial pipelines via Spark UI diagnosis: **data-skew remediation** (salting, broadcast joins), Adaptive Query Execution, caching, and partition-aware incremental loads.
- Built **PySpark/Spark SQL transformation jobs** (filtering, complex joins, aggregations, deduplication, enrichment) with **source-to-target validation** for regulatory compliance reporting.
- Performed **data quality validation, reconciliation, duplicate checks, and variance checks** across large financial datasets; maintained audit-ready documentation.
- Automated batch orchestration in **Apache Airflow** (retry logic, dependency management, failure alerting); maintained **CI/CD with Git/Jenkins** across DEV/QA/PROD.

- Converted **procedural ETL logic into set-based execution** patterns for cloud targets; registered data lake assets in **Glue Data Catalog** enabling metadata discovery across Spark, SQL, and Redshift.
- Enforced **data governance controls** on financial datasets: validation rules, lineage capture, and compliance documentation for regulatory audits.
- Delivered **knowledge-transfer sessions and operational runbooks** to client engineering teams for the modernized platform.

Experis IT (Thomson Reuters) — India

ETL Developer | June 2017 – Sep 2019

- Built and supported **large-scale extraction workflows** from legacy SQL Server databases into AWS S3 data lake; developed transformation jobs (Glue, PySpark, HiveQL) for sales, subscription, and revenue data.
- Wrote **complex SQL** (window functions, subqueries, multi-table joins) for data validation, ad-hoc analysis, and query tuning.
- Implemented **daily data validation**: missing-record detection, duplicate checks, source-to-target row-count validation, and revenue variance reviews.
- Tuned batch performance (~20–25% write-delay reduction) via Parquet output optimization and partition strategies.
- Monitored pipeline execution, **diagnosed failed loads, and troubleshoot ETL bottlenecks** to protect reporting SLAs; proactively resolved issues to **maximize uptime**; administered scheduled jobs via **SQL Server Agent** and cron.
- Developed **stored procedures, functions, and views (T-SQL)** supporting ETL and reporting workloads.
- Automated routine job operations with **UNIX shell scripting**; maintained ETL design documentation and source-to-target mapping records.

EDUCATION

Masters in Information Technology Management, 2024 — California Baptist University, Riverside, CA

B.Tech in Computer Science and Engineering, 2017 — JNTU-K University, India